

U-Net-architecture

March 5, 2023

```
[26]: import torch
import torchvision
from torchvision import models
import torchvision.transforms as T

import numpy as np
from PIL import Image
import matplotlib.pyplot as plt

print('pytorch', torch.__version__)
print('torchvision', torchvision.__version__)
```

pytorch 1.13.1
torchvision 0.14.1

```
[18]: path_data = Path('SegTMS/train/')
path_lbl = path_data/'Labels'
path_img = path_data/'Images'
```

```
[44]: IMG_SIZE = (256,256)

COLORS = np.array([
    (0, 0, 0),          # 0=background
    (128, 0, 0),        # 1=aeroplane
    (0, 128, 0),        # 2=bicycle
    (128, 128, 0),      # 3=bird
    (0, 0, 128),        # 4=boat
    (128, 0, 128),      # 5=bottle
    (0, 128, 128),      # 6=bus
    (128, 128, 128),    # 7=car
    (255, 255, 255),    # 8=cat
    (192, 0, 0),         # 9=chair
    (64, 128, 0),        # 10=cow
    (192, 128, 0),       # 11=dining table
    (64, 0, 128),        # 12=dog
    (192, 0, 128),       # 13=horse
    (64, 128, 128),      # 14=motorbike
    (192, 128, 128),     # 15=person
```

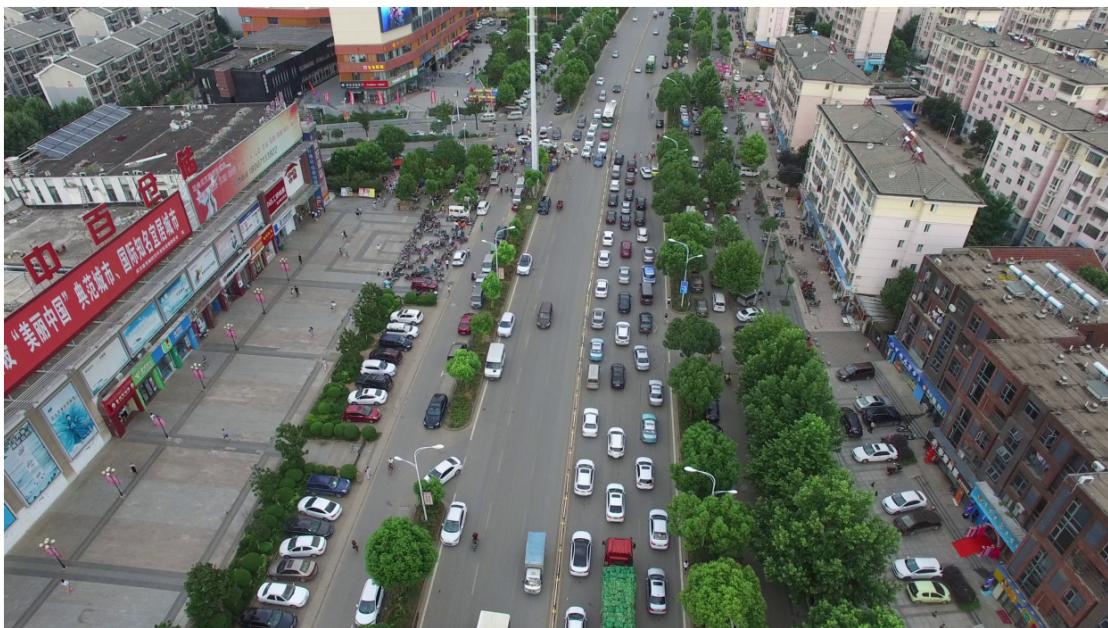
```
(0, 64, 0),      # 16=potted plant  
(128, 64, 0),   # 17=sheep  
(0, 192, 0),    # 18=sofa  
(128, 192, 0),  # 19=train  
(0, 64, 128)    # 20=tv/monitor  
])
```

[45]: *#Download Model*
deeplab = models.segmentation.deeplabv3_resnet101(pretrained=True).eval()

[46]: *#Load Image*
img = Image.open('SegTMS/train/Images/img3.png')

plt.figure(figsize=(16, 16))
plt.imshow(img)
plt.axis('off')

[46]: (-0.5, 3839.5, 2159.5, -0.5)



[52]: *#Image to Tensor*
trf = T.Compose([
 T.Resize(IMG_SIZE),
 T.CenterCrop(IMG_SIZE), *# make square image*
 T.ToTensor(),
 T.Normalize(
 mean=[0.485, 0.456, 0.406],

```
        std=[0.229, 0.224, 0.225]
    )
])

input_img = trf(img).unsqueeze(0)
```

```
[53]: out = deeplab(input_img) ['out']

print(out.shape)
```

```
torch.Size([1, 21, 256, 256])
```

```
[54]: #Extract Class Map
```

```
out = torch.argmax(out.squeeze(), dim=0)
out = out.detach().cpu().numpy()

print(out.shape)
print(np.unique(out))
```

```
(256, 256)
[0]
```

```
[55]: def seg_map(img, n_classes=21):
    rgb = np.zeros((img.shape[0], img.shape[1], 3), dtype=np.uint8)

    for c in range(n_classes):
        idx = img == c

        rgb[idx] = COLORS[c]

    return rgb
```

```
[88]: #Class Map to Segmentation Map
```

```
out_seg = seg_map(out)

fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(16, 16))
ax[0].imshow(img)
ax[0].axis('off')
ax[1].imshow(out_seg)
ax[1].axis('off')
```

```
[88]: (-0.5, 454.5, 255.5, -0.5)
```



```
[202]: ##-----Semantic Segmentation-Pytorch
      ↳unet-----#
```

```
[281]: from torch_snippets import *
from torchvision import transforms
from sklearn.model_selection import train_test_split
import cv2
import glob
from torchsummary import summary
import segmentation_models_pytorch as smp
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms as T
import torchvision
import torch.nn.functional as F
from torch.autograd import Variable
from torch.utils.data import DataLoader
from transformers import SegformerFeatureExtractor
from transformers import SegformerForSemanticSegmentation
device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
[282]: tfms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]) #↳
      ↳imagenet
])
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
[283]: #Batch generation
labels=sorted(glob.glob('SegTMS/train/Labels/*.png'))
images=sorted(glob.glob('SegTMS/train/Images/*.png'))
train_images=images[:int(0.8*(len(images)))]
train_labels=labels[:int(0.8*(len(labels)))]
test_images=images[int(0.8*(len(images))):]
```

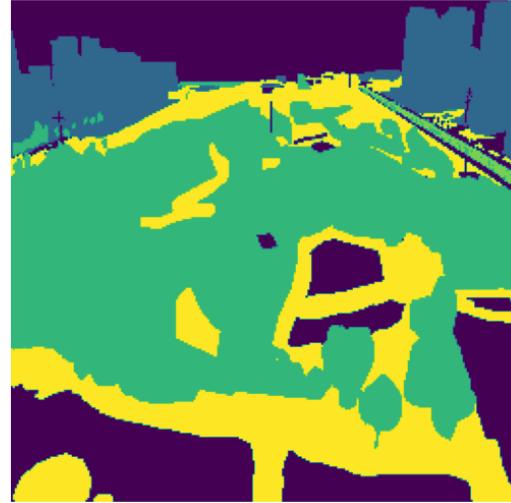
```
test_label=labels[int(0.8*(len(labels))):]
```

[284]: w=256
h=256

```
class SegData(Dataset):  
    def __init__(self,images,labels):  
        self.images=images  
        self.labels=labels  
    def __len__(self):  
        return len(self.images)  
    def __getitem__(self, ix):  
        image = read(self.images[ix], 1)  
        image = cv2.resize(image, (w,h))  
        mask = read(self.labels[ix],1)  
        mask = cv2.resize(mask, (w,h))  
        mask=cv2.cvtColor(mask,cv2.COLOR_RGB2GRAY)  
        return image, mask  
    def choose(self): return self[randint(len(self))]  
    def collate_fn(self, batch):  
        ims, masks = list(zip(*batch))  
        ims = torch.cat([tfms(im.copy()/255.)[None] for im in ims]).float().  
        to(device)  
        ce_masks = torch.cat([torch.tensor(mask[None]) for mask in masks]).  
        long().to(device)  
        return ims, ce_masks
```

[285]: train_dataset=SegData(train_images,train_labels)
test_dataset=SegData(test_images,test_label)

[286]: image,mask=train_dataset[-1]
plt.figure(figsize = (14, 10))
plt.subplot(1,2,1)
plt.axis('off')
plt.imshow(np.array(image))
plt.subplot(1,2,2)
plt.imshow(np.array(mask).squeeze())
plt.axis('off')
plt.show()
print(image.shape)
print(mask.shape)
print(mask.max())



(256, 256, 3)

(256, 256)

113

```
[287]: nbatch_size=16
train_loader = torch.utils.data.DataLoader(train_dataset, □
    ↪batch_size=nbatch_size, shuffle=True, num_workers=0)
test_loader=torch.utils.data.
    ↪DataLoader(test_dataset,batch_size=nbatch_size,shuffle=True, num_workers=0)
```

```
[288]: # lets look at some samples
image, mask = train_dataset[0]
plt.figure(figsize = (14, 10))
plt.subplot(1,2,1)
plt.axis('off')
plt.imshow(np.array(image))
plt.subplot(1,2,2)
plt.imshow(np.array(mask).squeeze())
plt.axis('off')
plt.show()

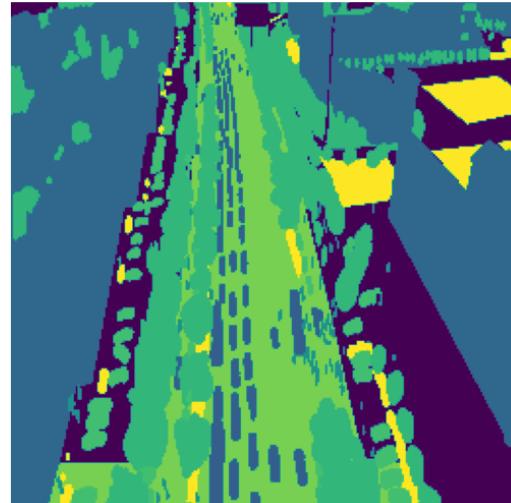
image, mask = test_dataset[0]

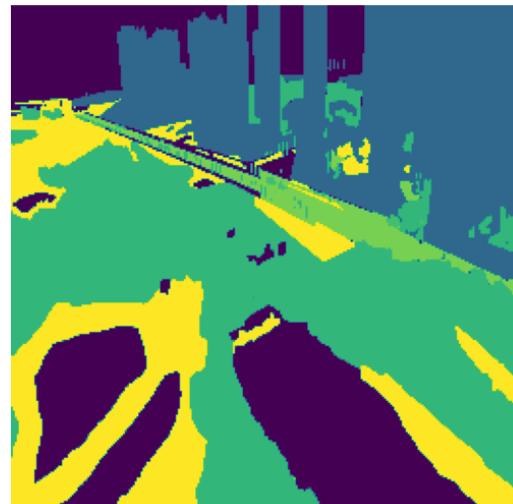
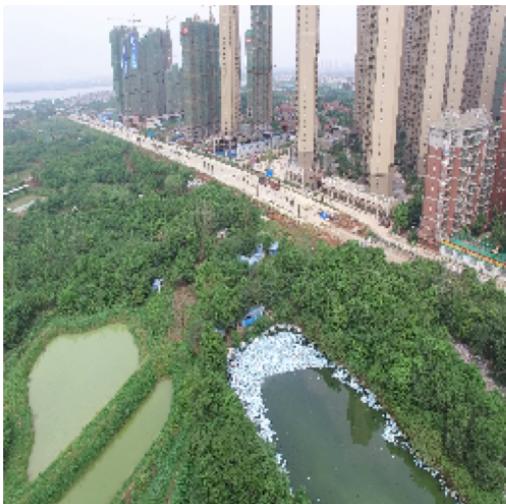
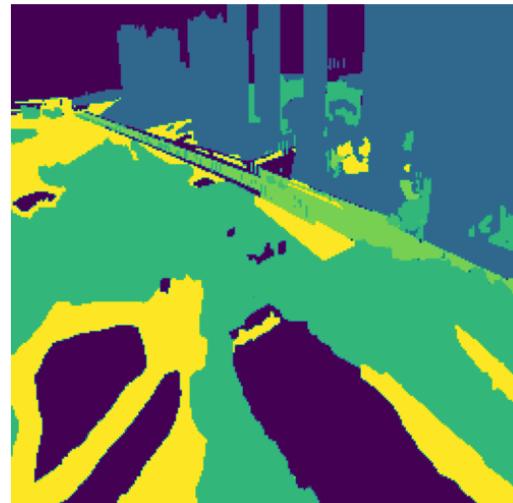
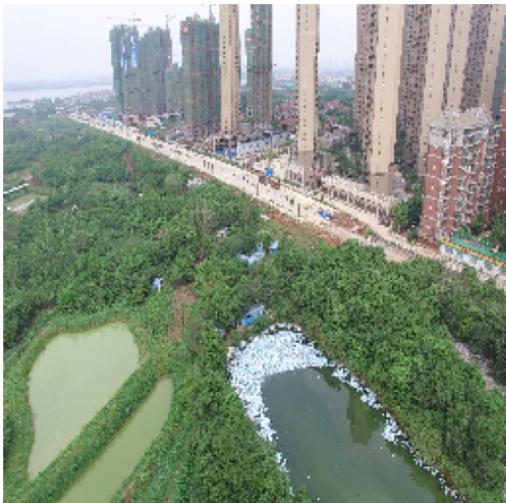
plt.figure(figsize = (14, 10))
plt.subplot(1,2,1)
plt.axis('off')
```

```
plt.imshow(np.array(image))
plt.subplot(1,2,2)
plt.imshow(np.array(mask).squeeze())
plt.axis('off')
plt.show()

image, mask = test_dataset[0]

plt.figure(figsize = (14, 10))
plt.subplot(1,2,1)
plt.imshow(np.array(image))
plt.axis('off')
plt.subplot(1,2,2)
plt.imshow(np.array(mask).squeeze())
plt.axis('off')
plt.show()
```





```
[293]: train_loader =train_loader  
val_loader=test_loader  
print('Train Size    : ', len(train_dataset))  
print('Val Size     : ', len(test_dataset))
```

Train Size : 160

Val Size : 40

```
[ ]: ##-----Semantic Segmentation-TF  
→unet-----#
```

```
[296]: import matplotlib.pyplot as plt
import imageio.v2 as imageio
import tensorflow as tf
```

```
import numpy as np

from tensorflow.keras.layers import Input
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import MaxPooling2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Conv2DTranspose
from tensorflow.keras.layers import concatenate
```

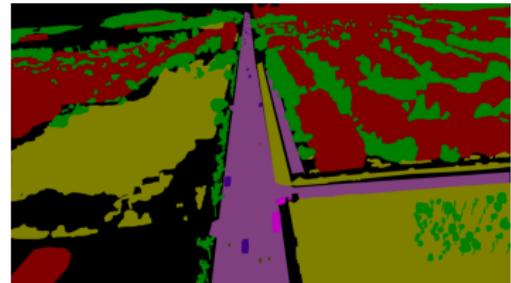
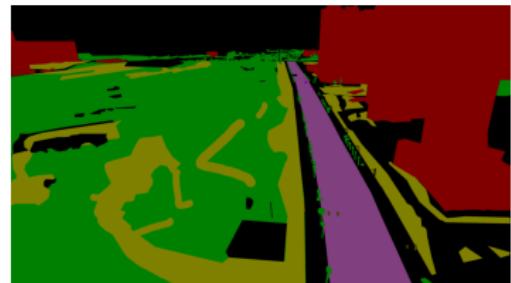
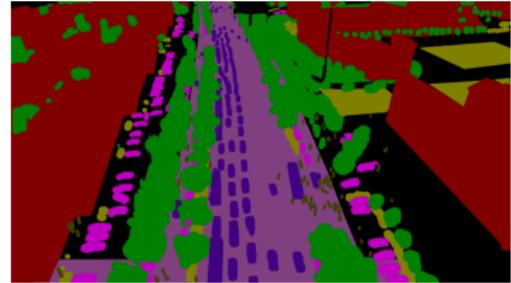
```
[308]: inp_dir = 'SegTMS/train/Images/'
mask_dir = 'SegTMS/train/Labels/'

inp_list = os.listdir(inp_dir)
mask_list= os.listdir(mask_dir)
inp_list.sort()
mask_list.sort()
inp_img_list = [os.path.join(inp_dir,i) for i in inp_list]
mask_list = [os.path.join(mask_dir,i) for i in mask_list]
```

```
[309]: def plt_image_and_mask_by_path(img_list,mask_lst,n=3):
    for i in range(3):
        img1=imageio.imread(img_list[i])
        mask1=imageio.imread(mask_lst[i])
        _,arr = plt.subplots(1,2,figsize=(10,10))
        arr[0].imshow(img1)
        arr[0].axis('off')
        arr[1].imshow(mask1)
        arr[1].axis('off')
```

```
[310]: def plt_image_and_mask_with_dataset(image_ds,mask_ds):
    _,arr = plt.subplots(1,2,figsize=(10,10))
    arr[0].imshow(tf.keras.preprocessing.image.array_to_img(image_ds))
    arr[0].axis('off')
    arr[1].imshow(tf.keras.preprocessing.image.array_to_img(mask_ds))
    arr[1].axis('off')
```

```
[312]: plt_image_and_mask_by_path(inp_img_list,mask_list,3)
```



```
[313]: # create images paths as tensors of file paths
image_names = tf.constant(inp_img_list)
mask_names = tf.constant(mask_list)

dataset = tf.data.Dataset.from_tensor_slices((image_names,mask_names))
for image,mask in dataset.take(1):
    print(image)
    print(mask)
```



```
tf.Tensor(b'SegTMS/train/Images/img1.png', shape=(), dtype=string)

tf.Tensor(b'SegTMS/train/Labels/mask1.png', shape=(), dtype=string)
```

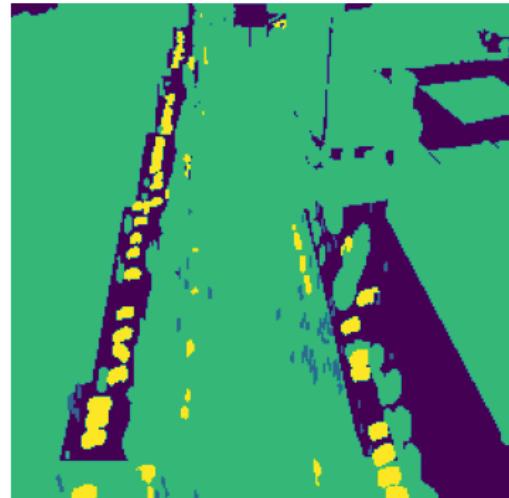
```
[328]: def process_path(image_,mask_):
    img = tf.io.read_file(image_)
    img = tf.image.decode_png(img,channels=3)
    img = tf.image.convert_image_dtype(img,tf.float32)

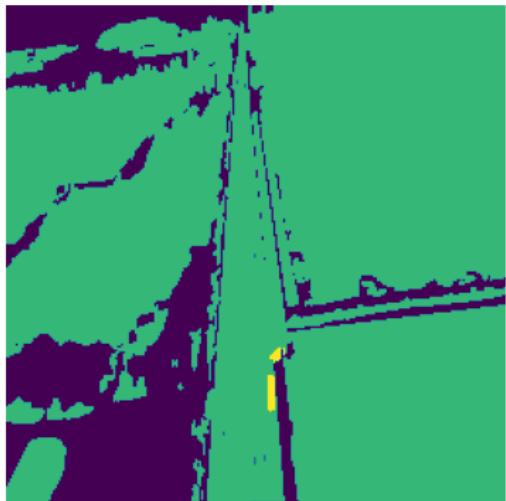
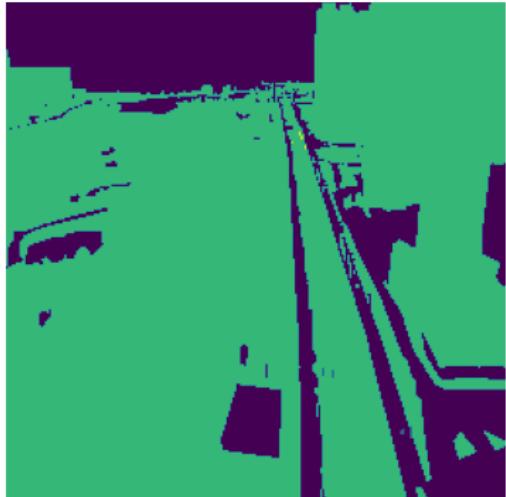
    mask = tf.io.read_file(mask_)
    mask = tf.image.decode_png(mask,channels=3)
    mask = tf.math.reduce_max(mask, axis=-1, keepdims=True)
    return img,mask

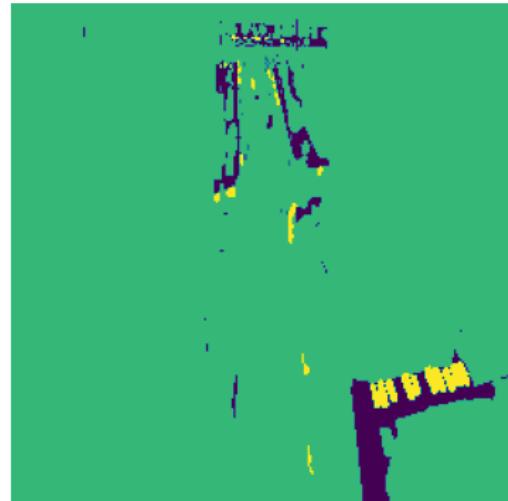
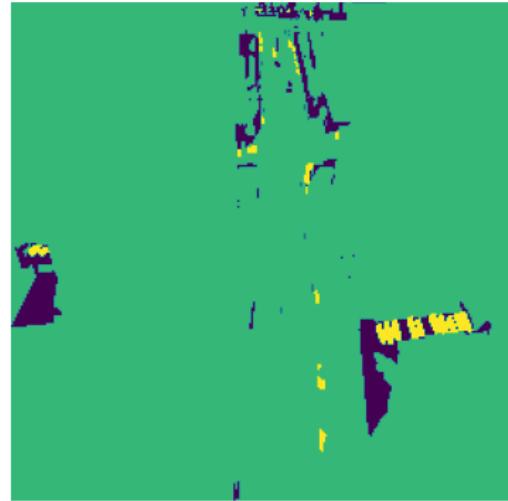
def process(image,mask):
    inp_image = tf.image.resize(image,(w,h),method='nearest')
    inp_mask = tf.image.resize(mask,(w,h),method='nearest')
    return inp_image,inp_mask

image_ds = dataset.map(process_path)
processed_images_ds = image_ds.map(process)
```

```
[330]: for image,mask in processed_images_ds.take(5):
    plt_image_and_mask_with_dataset(image,mask)
```







```
[331]: def conv_block(inputs=None, n_filters=64, dropout_prob=0, max_pooling=True):
    conv = Conv2D(n_filters, # Number of filters
                  kernel_size=3, # Kernel size
                  activation='relu',
                  padding='same',
                  kernel_initializer='he_normal')(inputs)
    conv = Conv2D(n_filters, # Number of filters
                  3, # Kernel size
                  activation='relu',
                  padding='same',
                  kernel_initializer='he_normal')(conv)
```

```

if dropout_prob > 0:
    conv = Dropout(dropout_prob)(conv)

if max_pooling:
    next_layer = MaxPooling2D(pool_size=(2,2))(conv)

else:
    next_layer = conv

skip_connection = conv

return next_layer, skip_connection

```

```

[332]: def upsampling_block(expansive_input, contractive_input, n_filters=64):
    up = Conv2DTranspose(
        filters=n_filters,      # number of filters
        kernel_size=(3,3),      # Kernel size
        strides=(2,2),
        padding='same')(expansive_input)

    merge = concatenate([up, contractive_input], axis=3)
    conv = Conv2D(filters=n_filters,      # Number of filters
                  kernel_size=(3,3),      # Kernel size
                  activation='relu',
                  padding='same',
                  kernel_initializer='he_normal')(merge)
    conv = Conv2D(n_filters,      # Number of filters
                  kernel_size=(3,3),      # Kernel size
                  activation='relu',
                  padding='same',
                  kernel_initializer='he_normal')(conv)

    return conv

```

```

[333]: def unet_model(input_size=(w, h, 3), n_filters=64, n_classes=256):

    inputs = Input(input_size)
    cblock1 = conv_block(inputs, n_filters)
    cblock2 = conv_block(cblock1[0], n_filters*2)
    cblock3 = conv_block(cblock2[0], n_filters*4)
    cblock4 = conv_block(cblock3[0], n_filters*8, dropout_prob=0.3)
    cblock5 = conv_block(cblock4[0], n_filters*16, dropout_prob=0.3, ↴
                         max_pooling=False)

    ublock6 = upsampling_block(cblock5[0], cblock4[1], n_filters*8)
    ublock7 = upsampling_block(ublock6, cblock3[1], n_filters*4)
    ublock8 = upsampling_block(ublock7, cblock2[1], n_filters*2)

```

```

ublock9 = upsampling_block(ublock8, cblock1[1], n_filters)

conv9 = Conv2D(n_filters,
            3,
            activation='relu',
            padding='same',
            kernel_initializer='he_normal')(ublock9)

conv10 = Conv2D(n_classes, 1, padding='same')(conv9)
model = tf.keras.Model(inputs=inputs, outputs=conv10)

return model

```

[334] : h=256
w=256
chs = 3
unet = unet_model((h, w, chs))

[335] : unet.summary()

```

Model: "model_3"
-----
-----  

Layer (type)           Output Shape        Param #  Connected to
-----  

-----  

input_4 (InputLayer)    [(None, 256, 256, 3  0
)]  

conv2d_60 (Conv2D)      (None, 256, 256, 64  1792
['input_4[0][0]'  

)  

conv2d_61 (Conv2D)      (None, 256, 256, 64  36928
['conv2d_60[0][0]'  

)  

max_pooling2d_12 (MaxPooling2D) (None, 128, 128, 64  0
['conv2d_61[0][0]'  

)  

conv2d_62 (Conv2D)      (None, 128, 128, 12  73856
['max_pooling2d_12[0][0]'  

8)  

conv2d_63 (Conv2D)      (None, 128, 128, 12  147584
['conv2d_62[0][0]'  

)
```

8)

```
max_pooling2d_13 (MaxPooling2D  (None, 64, 64, 128)  0
['conv2d_63[0][0]')
)

conv2d_64 (Conv2D)          (None, 64, 64, 256)  295168
['max_pooling2d_13[0][0]']

conv2d_65 (Conv2D)          (None, 64, 64, 256)  590080
['conv2d_64[0][0]']

max_pooling2d_14 (MaxPooling2D  (None, 32, 32, 256)  0
['conv2d_65[0][0]')
)

conv2d_66 (Conv2D)          (None, 32, 32, 512)  1180160
['max_pooling2d_14[0][0]']

conv2d_67 (Conv2D)          (None, 32, 32, 512)  2359808
['conv2d_66[0][0]']

dropout_6 (Dropout)         (None, 32, 32, 512)  0
['conv2d_67[0][0]']

max_pooling2d_15 (MaxPooling2D  (None, 16, 16, 512)  0
['dropout_6[0][0]')
)

conv2d_68 (Conv2D)          (None, 16, 16, 1024)  4719616
['max_pooling2d_15[0][0]')
)

conv2d_69 (Conv2D)          (None, 16, 16, 1024)  9438208
['conv2d_68[0][0]')
)

dropout_7 (Dropout)         (None, 16, 16, 1024)  0
['conv2d_69[0][0]')
)

conv2d_transpose_12 (Conv2DTra  (None, 32, 32, 512)  4719104
['dropout_7[0][0]')
nspose)

concatenate_12 (Concatenate) (None, 32, 32, 1024)  0
['conv2d_transpose_12[0][0]', )
)
```

```

'dropout_6[0][0]']

conv2d_70 (Conv2D)           (None, 32, 32, 512) 4719104
['concatenate_12[0][0]']

conv2d_71 (Conv2D)           (None, 32, 32, 512) 2359808
['conv2d_70[0][0]']

conv2d_transpose_13 (Conv2DTra (None, 64, 64, 256) 1179904
['conv2d_71[0][0]']
nspose)

concatenate_13 (Concatenate) (None, 64, 64, 512) 0
['conv2d_transpose_13[0][0]', 'conv2d_65[0][0]']

conv2d_72 (Conv2D)           (None, 64, 64, 256) 1179904
['concatenate_13[0][0]']

conv2d_73 (Conv2D)           (None, 64, 64, 256) 590080
['conv2d_72[0][0]']

conv2d_transpose_14 (Conv2DTra (None, 128, 128, 12 295040
['conv2d_73[0][0]']
nspose)                      8)

concatenate_14 (Concatenate) (None, 128, 128, 25 0
['conv2d_transpose_14[0][0]', 6)
'conv2d_63[0][0]']

conv2d_74 (Conv2D)           (None, 128, 128, 12 295040
['concatenate_14[0][0]']
8)

conv2d_75 (Conv2D)           (None, 128, 128, 12 147584
['conv2d_74[0][0]']
8)

conv2d_transpose_15 (Conv2DTra (None, 256, 256, 64 73792
['conv2d_75[0][0]']
nspose)                      )

concatenate_15 (Concatenate) (None, 256, 256, 12 0
['conv2d_transpose_15[0][0]', 8)
'conv2d_61[0][0]']

```

```

conv2d_76 (Conv2D)           (None, 256, 256, 64  73792
['concatenate_15[0][0]']
)
conv2d_77 (Conv2D)           (None, 256, 256, 64  36928
['conv2d_76[0][0]']
)
conv2d_78 (Conv2D)           (None, 256, 256, 64  36928
['conv2d_77[0][0]']
)
conv2d_79 (Conv2D)           (None, 256, 256, 25  16640
['conv2d_78[0][0]']
6)

=====
=====

Total params: 34,566,848
Trainable params: 34,566,848
Non-trainable params: 0
-----
```

```
[336]: unet.compile(optimizer='adam',
                    loss=tf.keras.losses.
                    SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

```
[ ]: EPOCHS = 10
VAL_SUBSPLITS = 5
BUFFER_SIZE = 500
BATCH_SIZE = 8
processed_images_ds.batch(BATCH_SIZE)
train_dataset = processed_images_ds.cache().shuffle(BUFFER_SIZE).
    batch(BATCH_SIZE)
print(processed_images_ds.element_spec)
model_history = unet.fit(train_dataset, epochs=EPOCHS)
```

```
(  
    TensorSpec(shape=(256, 256, 3), dtype=tf.float32, name=None),  
    TensorSpec(shape=(256, 256, 1), dtype=tf.uint8, name=None)  
)
```

Epoch 1/10

2023-03-05 04:26:24.133550: I
tensorflow/core/kernels/data/shuffle_dataset_op.cc:392] Filling up shuffle

```
buffer (this may take a while): 55 of 500
2023-03-05 04:26:34.459327: I
tensorflow/core/kernels/data/shuffle_dataset_op.cc:392] Filling up shuffle
buffer (this may take a while): 120 of 500
2023-03-05 04:26:44.536709: I
tensorflow/core/kernels/data/shuffle_dataset_op.cc:392] Filling up shuffle
buffer (this may take a while): 177 of 500
2023-03-05 04:26:48.136355: I
tensorflow/core/kernels/data/shuffle_dataset_op.cc:417] Shuffle buffer filled.

25/25 [=====] - 809s 30s/step - loss: 1.6286 -
accuracy: 0.6886
Epoch 2/10
25/25 [=====] - 743s 30s/step - loss: 0.6234 -
accuracy: 0.8106
Epoch 3/10
1/25 [>...] - ETA: 13:09 - loss: 0.6287 - accuracy:
0.7670
```

```
[ ]: unet.save('seg_surveillance')
model = tf.keras.models.load_model('seg_surveillance')
```

```
[ ]: acc = model_history.history['accuracy']
loss = model_history.history['loss']

epochs_range = range(100)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.legend(loc='lower right')
plt.title('Training Accuracy')
plt.show()

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, loss, label='Training Loss')
plt.legend(loc='lower right')
plt.title('Training loss')
plt.show()
```

```
[ ]: def display(display_list):
    plt.figure(figsize=(15, 15))
    title = ['Input Image', 'True Mask', 'Predicted Mask']
    for i in range(len(display_list)):
        plt.subplot(1, len(display_list), i+1)
        plt.title(title[i])
        plt.imshow(tf.keras.preprocessing.image.array_to_img(display_list[i]))
```

```
plt.axis('off')
plt.show()

[ ]: def create_mask(pred_mask):
    pred_mask = tf.argmax(pred_mask, axis=-1)
    pred_mask = pred_mask[..., tf.newaxis]
    return pred_mask[0]

[ ]: def show_predictions(dataset=None, num=1):
    if dataset:
        for image, mask in dataset.take(num):
            pred_mask = unet.predict(image)
            display([image[0], mask[0], create_mask(pred_mask)])
```

[]: show_predictions(train_dataset, 6)

[]:

[]: